

Introduction_to_SQL

1. Explain the fundamental differences between DDL, DML, and DQL commands in SQL. Provide one example for each type of command.

* DDL (Data Definition Language) is used to define or change the structure of database objects like tables.

Example:

```
CREATE TABLE Students (ID INT, Name VARCHAR(50));
```

* DML (Data Manipulation Language) is used to insert, update, or delete data from tables.

Example:

```
INSERT INTO Students VALUES (1, 'John');
```

* DQL (Data Query Language) is used to fetch data from the database.

Example:

```
SELECT * FROM Students;
```

2. What is the purpose of SQL constraints? Name and describe three common types of constraints, providing a simple scenario where each would be useful.

* SQL constraints are rules applied to table columns to ensure valid and consistent data is stored in the database.

* PRIMARY KEY: Uniquely identifies each row and does not allow NULL values.

Example: StudentID in a Students table.

* UNIQUE: Ensures all values in a column are different.

Example: Email column in a Users table must be unique.

* FOREIGN KEY: Creates a relationship between two tables.

Example: CustomerID in Orders table refers to CustomerID in Customers table.

3. Explain the difference between LIMIT and OFFSET clauses in SQL. How would you use them together to retrieve the third page of results, assuming each page has 10 records?

- * The LIMIT clause is used to restrict the number of rows returned in a result, while OFFSET skips a specified number of rows before starting to return rows.

- * To get the 3rd page of results with 10 records per page, you skip the first 20 records and show the next 10:

```
SELECT * FROM table_name  
LIMIT 10 OFFSET 20;
```

4. What is a Common Table Expression (CTE) in SQL, and what are its main benefits? Provide a simple SQL example demonstrating its usage.

- * A CTE (Common Table Expression) is a temporary result set that can be referenced within a SELECT, INSERT, UPDATE, or DELETE statement. It makes complex queries easier to read and manage.

Example:

```
WITH HighSalary AS (  
    SELECT name, salary FROM Employees WHERE salary > 50000  
)  
SELECT * FROM HighSalary;
```

- * This creates a temporary result named HighSalary and then selects from it.

5. Describe the concept of SQL Normalization and its primary goals. Briefly explain the first three normal forms (1NF, 2NF, 3NF).

- * Normalization in SQL is the process of organizing data to reduce redundancy and improve data integrity.

- * 1NF (First Normal Form): Ensures each column has atomic (single) values and no repeating groups.

- * 2NF (Second Normal Form): Must be in 1NF and all non-key columns must depend on the whole primary key (applies to composite keys).
- * 3NF (Third Normal Form): Must be in 2NF and no non-key column should depend on another non-key column.

6. Create a database named ECommerceDB and perform the following tasks:

1. Create the following tables with appropriate data types and constraints:

-- Create Database (skip this part if using an online editor)

```
CREATE DATABASE ECommerceDB;
```

```
USE ECommerceDB;
```

-- Create Categories table

```
CREATE TABLE Categories (  
    CategoryID INT PRIMARY KEY,  
    CategoryName VARCHAR(50) NOT NULL UNIQUE  
);
```

-- Create Products table

```
CREATE TABLE Products (  
    ProductID INT PRIMARY KEY,  
    ProductName VARCHAR(100) NOT NULL UNIQUE,  
    CategoryID INT,  
    Price DECIMAL(10,2) NOT NULL,  
    StockQuantity INT,  
    FOREIGN KEY (CategoryID) REFERENCES Categories(CategoryID)  
);
```

```
-- Create Customers table
CREATE TABLE Customers (
  CustomerID INT PRIMARY KEY,
  CustomerName VARCHAR(100) NOT NULL,
  Email VARCHAR(100) UNIQUE,
  JoinDate DATE
);


-- Create Orders table
CREATE TABLE Orders (
  OrderID INT PRIMARY KEY,
  CustomerID INT,
  OrderDate DATE NOT NULL,
  TotalAmount DECIMAL(10,2),
  FOREIGN KEY (CustomerID) REFERENCES Customers(CustomerID)
);


SHOW TABLES;

DESCRIBE Categories;
DESCRIBE Products;
DESCRIBE Customers;
DESCRIBE Orders;
```

Result Grid		Filter Rows:
	Tables_in_ecommerceadb	
▶	categories	
	customers	
	orders	
	products	


Result Grid



Filter Rows:

Export:

Wrap Cell Con

	Field	Type	Null	Key	Default	Extra
▶	ProductID	int	NO	PRI	NULL	
	ProductName	varchar(100)	NO	UNI	NULL	
	CategoryID	int	YES	MUL	NULL	
	Price	decimal(10,2)	NO		NULL	
	StockQuantity	int	YES		NULL	

Result Grid


Filter Rows:

Export:


Wrap Cell

	Field	Type	Null	Key	Default	Extra
▶	CustomerID	int	NO	PRI	NULL	
	CustomerName	varchar(100)	NO		NULL	
	Email	varchar(100)	YES	UNI	NULL	
	JoinDate	date	YES		NULL	

Result Grid

Filter Rows:

Export:

Wrap C

	Field	Type	Null	Key	Default	Extra
▶	OrderID	int	NO	PRI	NULL	
	CustomerID	int	YES	MUL	NULL	
	OrderDate	date	NO		NULL	
	TotalAmount	decimal(10,2)	YES		NULL	

Result Grid		Filter Rows:	Export:		Wrap	
	Field	Type	Null	Key	Default	Extra
▶	OrderID	int	NO	PRI	NULL	
	CustomerID	int	YES	MUL	NULL	
	OrderDate	date	NO		NULL	
	TotalAmount	decimal(10,2)	YES		NULL	

2. Insert the following records into each table

INSERT INTO Categories (CategoryID, CategoryName) VALUES

(1, 'Electronics'),

(2, 'Books'),

(3, 'Home Goods'),

(4, 'Apparel');

INSERT INTO Customers (CustomerID, CustomerName, Email, JoinDate)
VALUES

(1, 'Alice Wonderland', 'alice@example.com', '2023-01-10'),

(2, 'Bob the Builder', 'bob@example.com', '2022-11-25'),

(3, 'Charlie Chaplin', 'charlie@example.com', '2023-03-01'),

(4, 'Diana Prince', 'diana@example.com', '2021-04-26');

INSERT INTO Products (ProductID, ProductName, CategoryID, Price,
StockQuantity) VALUES

(101, 'Laptop Pro', 1, 1200.00, 50),

(102, 'SQL Handbook', 2, 45.50, 200),

(103, 'Smart Speaker', 1, 99.99, 150),

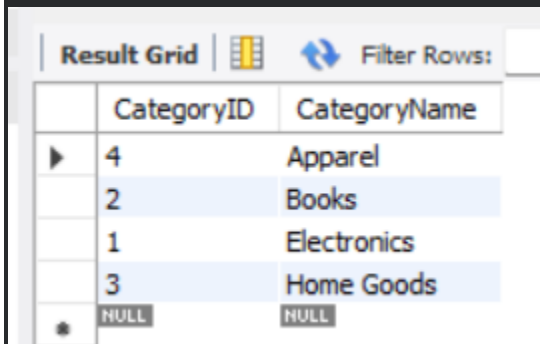
(104, 'Coffee Maker', 3, 75.00, 80),

(105, 'Novel : The Great SQL', 2, 25.00, 120),

```
(106, 'Wireless Earbuds', 1, 150.00, 100),  
(107, 'Blender X', 3, 120.00, 60),  
(108, 'T-Shirt Casual', 4, 20.00, 300);
```

```
INSERT INTO Orders (OrderID, CustomerID, OrderDate, TotalAmount) VALUES  
(1001, 1, '2023-04-26', 1245.50),  
(1002, 2, '2023-10-12', 99.99),  
(1003, 1, '2023-07-01', 145.00),  
(1004, 3, '2023-01-14', 150.00),  
(1005, 2, '2023-09-24', 120.00),  
(1006, 1, '2023-06-19', 20.00);
```

```
SELECT * FROM Categories;  
SELECT * FROM Products;  
SELECT * FROM Customers;  
SELECT * FROM Orders;
```



The screenshot shows a database interface with a 'Result Grid' and a 'Filter Rows' button. The grid displays a table with two columns: 'CategoryID' and 'CategoryName'. The data is as follows:

	CategoryID	CategoryName
▶	4	Apparel
	2	Books
	1	Electronics
	3	Home Goods
•	NULL	NULL

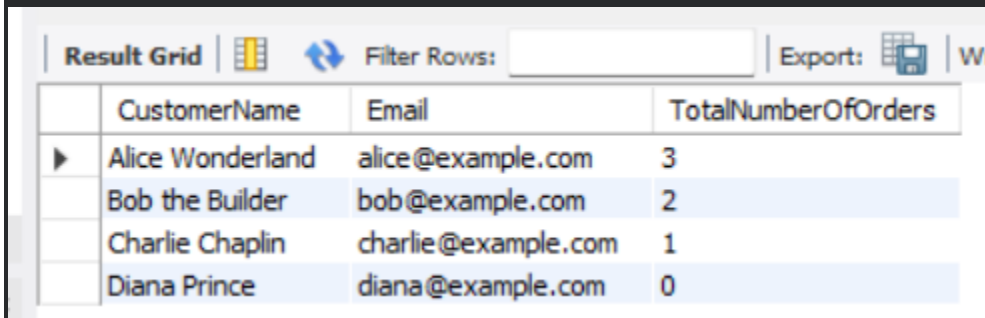
Result Grid Filter Rows: <input type="text"/> Edit: Export/					
	ProductID	ProductName	CategoryID	Price	StockQuantity
▶	101	Laptop Pro	1	1200.00	50
	102	SQL Handbook	2	45.50	200
	103	Smart Speaker	1	99.99	150
	104	Coffee Maker	3	75.00	80
	105	Novel : The Great SQL	2	25.00	120
	106	Wireless Earbuds	1	150.00	100
	107	Blender X	3	120.00	60
	108	T-Shirt Casual	4	20.00	300
*	NULL	NULL	NULL	NULL	NULL

Result Grid Filter Rows: <input type="text"/> Edit:				
	CustomerID	CustomerName	Email	JoinDate
▶	1	Alice Wonderland	alice@example.com	2023-01-10
	2	Bob the Builder	bob@example.com	2022-11-25
	3	Charlie Chaplin	charlie@example.com	2023-03-01
	4	Diana Prince	diana@example.com	2021-04-26
*	NULL	NULL	NULL	NULL

Result Grid Filter Rows: <input type="text"/> Edit:				
	OrderID	CustomerID	OrderDate	TotalAmount
▶	1001	1	2023-04-26	1245.50
	1002	2	2023-10-12	99.99
	1003	1	2023-07-01	145.00
	1004	3	2023-01-14	150.00
	1005	2	2023-09-24	120.00
	1006	1	2023-06-19	20.00
*	NULL	NULL	NULL	NULL

7. Generate a report showing CustomerName, Email, and the TotalNumberOfOrders for each customer. Include customers who have not placed any orders, in which case their TotalNumberOfOrders should be 0. Order the results by CustomerName.


```
SELECT
    c.CustomerName,
    c.Email,
    COUNT(o.OrderID) AS TotalNumberOfOrders
FROM
    Customers c
LEFT JOIN
    Orders o
    ON c.CustomerID = o.CustomerID
GROUP BY
    c.CustomerID, c.CustomerName, c.Email
ORDER BY
    c.CustomerName ASC;
```

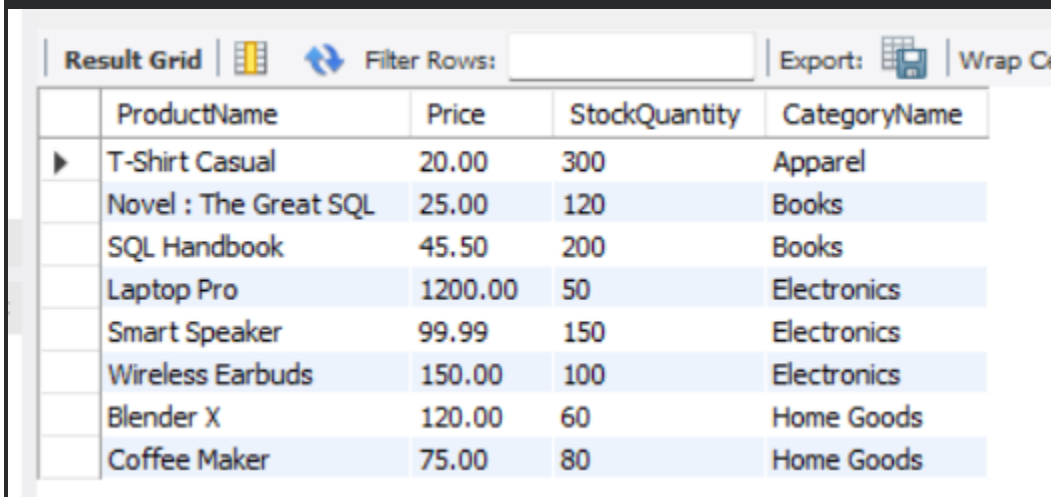


The screenshot shows a database application window with a 'Result Grid' tab. The grid displays the results of the SQL query, sorted by CustomerName in ascending order. The columns are CustomerName, Email, and TotalNumberOfOrders. The data rows are: Alice Wonderland (3 orders), Bob the Builder (2 orders), Charlie Chaplin (1 order), and Diana Prince (0 orders). The interface includes a 'Filter Rows' search bar and an 'Export' button.

	CustomerName	Email	TotalNumberOfOrders
▶	Alice Wonderland	alice@example.com	3
	Bob the Builder	bob@example.com	2
	Charlie Chaplin	charlie@example.com	1
	Diana Prince	diana@example.com	0

8. Retrieve Product Information with Category: Write a SQL query to display the ProductName, Price, StockQuantity, and CategoryName for all products. Order the results by CategoryName and then ProductName alphabetically.

```
SELECT
    p.ProductName,
    p.Price,
    p.StockQuantity,
    c.CategoryName
FROM
    Products p
INNER JOIN
    Categories c
    ON p.CategoryID = c.CategoryID
ORDER BY
    c.CategoryName ASC,
    p.ProductName ASC;
```

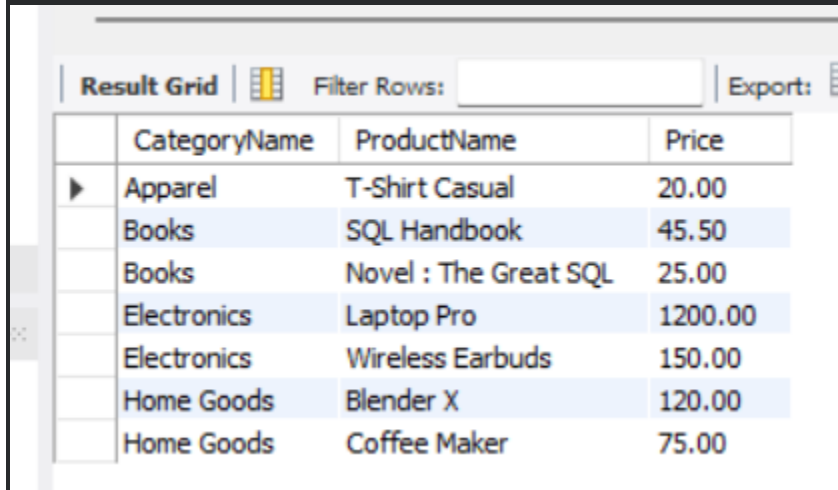


	ProductName	Price	StockQuantity	CategoryName
▶	T-Shirt Casual	20.00	300	Apparel
	Novel : The Great SQL	25.00	120	Books
	SQL Handbook	45.50	200	Books
	Laptop Pro	1200.00	50	Electronics
	Smart Speaker	99.99	150	Electronics
	Wireless Earbuds	150.00	100	Electronics
	Blender X	120.00	60	Home Goods
	Coffee Maker	75.00	80	Home Goods

9. Write a SQL query that uses a Common Table Expression (CTE) and a Window Function (specifically ROW_NUMBER() or RANK()) to display the CategoryName, ProductName, and Price for the top 2 most expensive products in each CategoryName.

```
WITH RankedProducts AS (  
    SELECT  
        c.CategoryName,  
        p.ProductName,  
        p.Price,  
        ROW_NUMBER() OVER (  
            PARTITION BY c.CategoryName  
            ORDER BY p.Price DESC  
        ) AS RankNo  
    FROM  
        Products p  
    INNER JOIN  
        Categories c  
        ON p.CategoryID = c.CategoryID  
)  
SELECT  
    CategoryName,  
    ProductName,  
    Price  
FROM  
    RankedProducts  
WHERE  
    RankNo <= 2  
ORDER BY
```

CategoryName ASC,
Price DESC;



The screenshot shows a 'Result Grid' window with a table of products. The table has four columns: an index column, 'CategoryName', 'ProductName', and 'Price'. The data is sorted by 'CategoryName' in ascending order and 'Price' in descending order. The categories shown are Apparel, Books, Electronics, and Home Goods. The 'Apparel' category has one item, 'T-Shirt Casual' for 20.00. The 'Books' category has two items: 'SQL Handbook' for 45.50 and 'Novel : The Great SQL' for 25.00. The 'Electronics' category has two items: 'Laptop Pro' for 1200.00 and 'Wireless Earbuds' for 150.00. The 'Home Goods' category has two items: 'Blender X' for 120.00 and 'Coffee Maker' for 75.00. The window also includes a 'Filter Rows' field and an 'Export' button.

	CategoryName	ProductName	Price
▶	Apparel	T-Shirt Casual	20.00
	Books	SQL Handbook	45.50
	Books	Novel : The Great SQL	25.00
	Electronics	Laptop Pro	1200.00
	Electronics	Wireless Earbuds	150.00
	Home Goods	Blender X	120.00
	Home Goods	Coffee Maker	75.00

Question 10 : You are hired as a data analyst by Sakila Video Rentals, a global movie rental company. The management team is looking to improve decision-making by analyzing existing customer, rental, and inventory data.

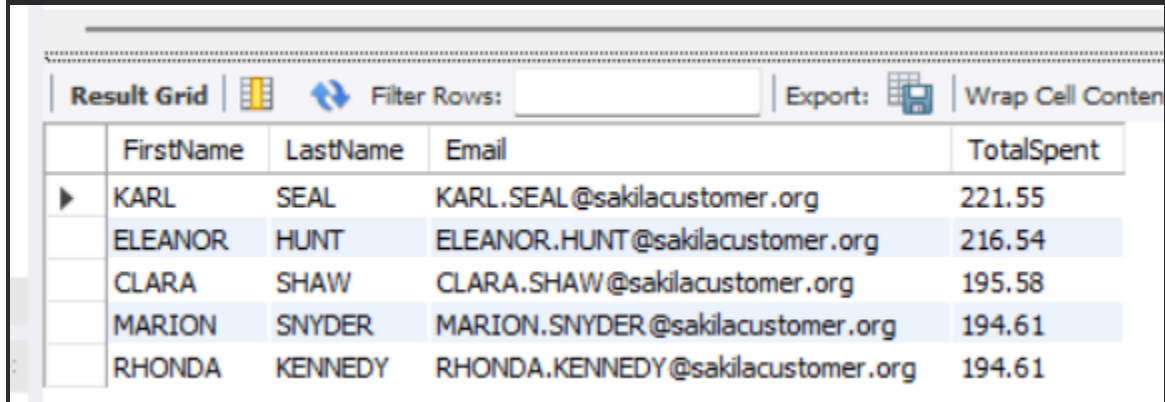
Identify the top 5 customers based on the total amount they've spent. Include customer name, email, and total amount spent.

```
SELECT
    c.first_name AS FirstName,
    c.last_name AS LastName,
    c.email AS Email,
    SUM(p.amount) AS TotalSpent
FROM
    customer c
JOIN
    payment p ON c.customer_id = p.customer_id
GROUP BY
    c.customer_id, c.first_name, c.last_name, c.email
```

ORDER BY

TotalSpent DESC

LIMIT 5;



The screenshot shows a database query result grid with a toolbar at the top. The toolbar includes a 'Result Grid' button, a 'Filter Rows' input field, an 'Export' button, and a 'Wrap Cell Content' button. The table below has five columns: 'FirstName', 'LastName', 'Email', and 'TotalSpent'. The data is sorted by 'TotalSpent' in descending order, showing the top 5 customers.

	FirstName	LastName	Email	TotalSpent
▶	KARL	SEAL	KARL.SEAL@sakilacustomer.org	221.55
	ELEANOR	HUNT	ELEANOR.HUNT@sakilacustomer.org	216.54
	CLARA	SHAW	CLARA.SHAW@sakilacustomer.org	195.58
	MARION	SNYDER	MARION.SNYDER@sakilacustomer.org	194.61
	RHONDA	KENNEDY	RHONDA.KENNEDY@sakilacustomer.org	194.61

Find the top 3 movie categories with the highest rental counts. Display the category name and number of rentals.

SELECT

cat.name AS CategoryName,

COUNT(r.rental_id) AS RentalCount

FROM

rental r

JOIN

inventory i ON r.inventory_id = i.inventory_id

JOIN

film f ON i.film_id = f.film_id

JOIN

film_category fc ON f.film_id = fc.film_id

JOIN

category cat ON fc.category_id = cat.category_id

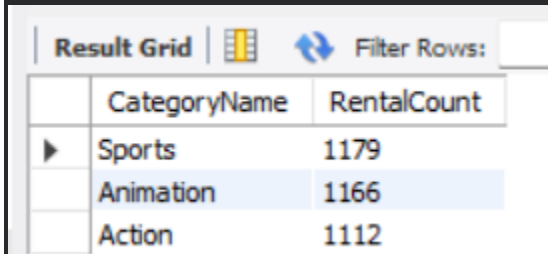
GROUP BY

cat.name

ORDER BY

RentalCount DESC

LIMIT 3;



A screenshot of a database application's 'Result Grid'. The grid has two columns: 'CategoryName' and 'RentalCount'. It displays three rows of data. The first row is 'Sports' with a rental count of 1179. The second row is 'Animation' with a rental count of 1166. The third row is 'Action' with a rental count of 1112. Above the grid, there are icons for 'Result Grid', a grid icon, and a 'Filter Rows' button with a search input field.

	CategoryName	RentalCount
▶	Sports	1179
	Animation	1166
	Action	1112

Calculate how many films are available at each store and how many of those have never been rented.

SELECT

s.store_id,

COUNT(i.inventory_id) AS TotalFilms,

SUM(CASE WHEN r.rental_id IS NULL THEN 1 ELSE 0 END) AS

NeverRented

FROM

store s

JOIN

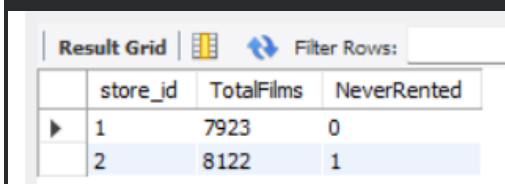
inventory i ON s.store_id = i.store_id

LEFT JOIN

rental r ON i.inventory_id = r.inventory_id

GROUP BY

s.store_id;

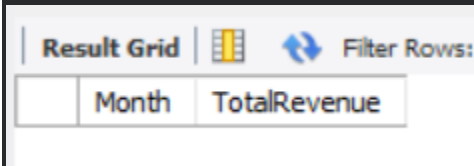


A screenshot of a database application's 'Result Grid'. The grid has four columns: 'store_id', 'TotalFilms', and 'NeverRented'. It displays two rows of data. The first row is for store_id 1, with 7923 total films and 0 never rented. The second row is for store_id 2, with 8122 total films and 1 never rented. Above the grid, there are icons for 'Result Grid', a grid icon, and a 'Filter Rows' button with a search input field.

	store_id	TotalFilms	NeverRented
▶	1	7923	0
	2	8122	1

Show the total revenue per month for the year 2023 to analyze business seasonality.

```
SELECT
    DATE_FORMAT(p.payment_date, '%Y-%m') AS Month,
    SUM(p.amount) AS TotalRevenue
FROM
    payment p
WHERE
    YEAR(p.payment_date) = 2023
GROUP BY
    DATE_FORMAT(p.payment_date, '%Y-%m')
ORDER BY
    Month;
```



The screenshot shows a database interface with a 'Result Grid' tab. Below the tab, there is a table with two columns: 'Month' and 'TotalRevenue'. The 'Month' column is currently selected.




Month	TotalRevenue
-------	--------------

Identify customers who have rented more than 10 times in the last 6 months.

```
SELECT
    c.first_name AS FirstName,
    c.last_name AS LastName,
    c.email AS Email,
    COUNT(r.rental_id) AS RentalsInLast6Months
FROM
    customer c
JOIN
    rental r ON c.customer_id = r.customer_id
WHERE
```

```
    r.rental_date >= DATE_SUB(CURDATE(), INTERVAL 6 MONTH)
GROUP BY
    c.customer_id, c.first_name, c.last_name, c.email
HAVING
    COUNT(r.rental_id) > 10
ORDER BY
    RentalsInLast6Months DESC;
```

Result Grid



Filter Rows:
Export:


	FirstName	LastName	Email	RentalsInLast6Months
--	-----------	----------	-------	----------------------